

INTRODUCTION TO DATA SCIENCE

Group 14 -- Project 1 -- Phase 1

Name	Student ID
Pham, Quoc Huy	2299356
Hussain, Zakiuddin	2338350
Lee, Daeul Haven	2308018

Link to colab: https://colab.research.google.com/drive/1zeUCdIEciMGKir-OK0m8sxObEmXK_QMF?usp=sharing

Link to the dataset: <https://archive.ics.uci.edu/dataset/791/metropt+3+dataset>

```
#Define ENV
USE_GG_DRIVE = 0

#Seting google Drive
if USE_GG_DRIVE:
    from goolge.colab import drive
    drive.mount('content\drive')

import pandas as pd
import seaborn as sns
import numpy as np

data = pd.read_csv("metro_data.csv")
print(data.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1516948 entries, 0 to 1516947
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Unnamed: 0             1516948 non-null int64   
1   timestamp              1516948 non-null object
2   TP2                    1516948 non-null float64
3   TP3                    1516948 non-null float64
4   H1                     1516948 non-null float64
5   DV_pressure            1516948 non-null float64
6   Reservoirs             1516948 non-null float64
7   Oil_temperature        1516948 non-null float64
8   Motor_current          1516948 non-null float64
9   COMP                   1516948 non-null float64
10  DV_eletric              1516948 non-null float64
11  Towers                  1516948 non-null float64
12  MPG                     1516948 non-null float64
13  LPS                     1516948 non-null float64
14  Pressure_switch         1516948 non-null float64
15  Oil_level               1516948 non-null float64
16  Caudal_impulses        1516948 non-null float64
dtypes: float64(15), int64(1), object(1)
memory usage: 196.7+ MB
None
```

I. Data Introduction

The MetroPT-3 dataset was created to support the development of predictive maintenance, anomaly detection, and remaining useful life (RUL) prediction models for compressors using deep learning and machine learning methods. It consists of multivariate time series data from several analogue and digital sensors installed on a train compressor. The data spans between February and August 2020 and includes 15 signals, such as pressures, motor current, oil temperature, and electrical signals of air intake valves. The dataset is eligible for incremental training and does not contain sensitive data. Data preprocessing includes data segmentation, normalization, and feature extraction. The dataset is unlabeled, but failure reports provided by the company are available for evaluating the effectiveness of anomaly detection, failure prediction, and RUL estimation algorithms. The dataset does not have missing values.

This dataset was chosen because it contains information that was obtained from actual events that accurately reflect real-world situations. This makes it a valuable resource for developing and testing algorithms in the field of anomaly detection, failure prediction, and remaining useful life estimation. Additionally, the availability of failure reports from the company allows us to validate the performance of our algorithms against real-world incidents. Also, this dataset is also represent time series attribute, which is crucial for analyzing trends and patterns over a specific period. The inclusion of time series data enables researchers to observe the progression of anomalies and failures, providing deeper insights into their causes and potential mitigation strategies.

II. Data Cleaning and Pre-processing

1) Data Overview

```
print(data.describe().round(2))
print(data.columns)
```

	Unnamed: 0	TP2	TP3	H1	DV_pressure	\
count	1516948.00	1516948.00	1516948.00	1516948.00	1516948.00	
mean	7584735.00	1.37	8.98	7.57	0.06	
std	4379053.12	3.25	0.64	3.33	0.38	
min	0.00	-0.03	0.73	-0.04	-0.03	
25%	3792367.50	-0.01	8.49	8.25	-0.02	
50%	7584735.00	-0.01	8.96	8.78	-0.02	
75%	11377102.50	-0.01	9.49	9.37	-0.02	
max	15169470.00	10.68	10.30	10.29	9.84	

	Reservoirs	Oil_temperature	Motor_current	COMP	DV_eletric	\
count	1516948.00	1516948.00	1516948.00	1516948.00	1516948.00	
mean	8.99	62.64	2.05	0.84	0.16	
std	0.64	6.52	2.30	0.37	0.37	
min	0.71	15.40	0.02	0.00	0.00	
25%	8.49	57.78	0.04	1.00	0.00	
50%	8.96	62.70	0.04	1.00	0.00	
75%	9.49	67.25	3.81	1.00	0.00	
max	10.30	89.05	9.30	1.00	1.00	

	Towers	MPG	LPS	Pressure_switch	Oil_level	\
count	1516948.00	1516948.00	1516948.00	1516948.00	1516948.00	
mean	0.92	0.83	0.00	0.99	0.90	
std	0.27	0.37	0.06	0.09	0.29	
min	0.00	0.00	0.00	0.00	0.00	
25%	1.00	1.00	0.00	1.00	1.00	
50%	1.00	1.00	0.00	1.00	1.00	
75%	1.00	1.00	0.00	1.00	1.00	
max	1.00	1.00	1.00	1.00	1.00	

	Caudal_impulses
count	1516948.00
mean	0.94
std	0.24
min	0.00
25%	1.00
50%	1.00
75%	1.00
max	1.00

```
Index(['Unnamed: 0', 'timestamp', 'TP2', 'TP3', 'H1', 'DV_pressure',
      'Reservoirs', 'Oil_temperature', 'Motor_current', 'COMP', 'DV_eletric',
      'Towers', 'MPG', 'LPS', 'Pressure_switch', 'Oil_level',
      'Caudal_impulses'],
      dtype='object')
```



```
print(data.head(10))
```

	Unnamed: 0	timestamp	TP2	TP3	H1	DV_pressure	\
0	0	2020-02-01 00:00:00	-0.012	9.358	9.340	-0.024	
1	10	2020-02-01 00:00:10	-0.014	9.348	9.332	-0.022	
2	20	2020-02-01 00:00:19	-0.012	9.338	9.322	-0.022	
3	30	2020-02-01 00:00:29	-0.012	9.328	9.312	-0.022	
4	40	2020-02-01 00:00:39	-0.012	9.318	9.302	-0.022	
5	50	2020-02-01 00:00:49	-0.012	9.306	9.290	-0.024	
6	60	2020-02-01 00:00:59	-0.012	9.296	9.280	-0.024	
7	70	2020-02-01 00:01:09	-0.014	9.286	9.270	-0.024	
8	80	2020-02-01 00:01:19	-0.012	9.276	9.258	-0.022	
9	90	2020-02-01 00:01:29	-0.012	9.264	9.248	-0.022	

	Reservoirs	Oil_temperature	Motor_current	COMP	DV_eletric	Towers	MPG	\
0	9.358	53.600	0.0400	1.0	0.0	1.0	1.0	

1	9.348	53.675	0.0400	1.0	0.0	1.0	1.0
2	9.338	53.600	0.0425	1.0	0.0	1.0	1.0
3	9.328	53.425	0.0400	1.0	0.0	1.0	1.0
4	9.318	53.475	0.0400	1.0	0.0	1.0	1.0
5	9.308	53.500	0.0400	1.0	0.0	1.0	1.0
6	9.298	53.375	0.0400	1.0	0.0	1.0	1.0
7	9.286	53.550	0.0400	1.0	0.0	1.0	1.0
8	9.276	53.425	0.0400	1.0	0.0	1.0	1.0
9	9.264	53.375	0.0400	1.0	0.0	1.0	1.0

	LPS	Pressure_switch	Oil_level	Caudal_impulses
0	0.0	1.0	1.0	1.0
1	0.0	1.0	1.0	1.0
2	0.0	1.0	1.0	1.0
3	0.0	1.0	1.0	1.0
4	0.0	1.0	1.0	1.0
5	0.0	1.0	1.0	1.0
6	0.0	1.0	1.0	1.0
7	0.0	1.0	1.0	1.0
8	0.0	1.0	1.0	1.0
9	0.0	1.0	1.0	1.0

2) Drop unnecessary columns

```
print(data.columns)
data = data.drop('Unnamed: 0', axis= 1)

Index(['Unnamed: 0', 'timestamp', 'TP2', 'TP3', 'H1', 'DV_pressure',
      'Reservoirs', 'Oil_temperature', 'Motor_current', 'COMP', 'DV_eletric',
      'Towers', 'MPG', 'LPS', 'Pressure_switch', 'Oil_level',
      'Caudal_impulses'],
      dtype='object')
```

3) Convert the timestamp collumn into pandas.DateTime data type

```
import datetime

#Check the current type of timestamp
print(f"Current type of timestamp is {type(data.timestamp[0])}")

#Convert timestamp to pandas.DateTime
data['timestamp'] = data['timestamp'].apply(pd.to_datetime, format = "%Y-%m-%d %H:%M:%S")

#Re-check the type
print(f"Current type of timestamp is {type(data.timestamp[0])}")

Current type of timestamp is <class 'str'>
Current type of timestamp is <class 'pandas._libs.tslibs.timestamps.Timestamp'>

print(data.head(10))
```

	timestamp	TP2	TP3	H1	DV_pressure	Reservoirs	\
0	2020-02-01 00:00:00	-0.012	9.358	9.340	-0.024	9.358	
1	2020-02-01 00:00:10	-0.014	9.348	9.332	-0.022	9.348	
2	2020-02-01 00:00:19	-0.012	9.338	9.322	-0.022	9.338	
3	2020-02-01 00:00:29	-0.012	9.328	9.312	-0.022	9.328	
4	2020-02-01 00:00:39	-0.012	9.318	9.302	-0.022	9.318	
5	2020-02-01 00:00:49	-0.012	9.306	9.290	-0.024	9.308	
6	2020-02-01 00:00:59	-0.012	9.296	9.280	-0.024	9.298	
7	2020-02-01 00:01:09	-0.014	9.286	9.270	-0.024	9.286	
8	2020-02-01 00:01:19	-0.012	9.276	9.258	-0.022	9.276	
9	2020-02-01 00:01:29	-0.012	9.264	9.248	-0.022	9.264	

	Oil_temperature	Motor_current	COMP	DV_eletric	Towers	MPG	LPS	\
0	53.600	0.0400	1.0	0.0	1.0	1.0	0.0	
1	53.675	0.0400	1.0	0.0	1.0	1.0	0.0	
2	53.600	0.0425	1.0	0.0	1.0	1.0	0.0	
3	53.425	0.0400	1.0	0.0	1.0	1.0	0.0	
4	53.475	0.0400	1.0	0.0	1.0	1.0	0.0	
5	53.500	0.0400	1.0	0.0	1.0	1.0	0.0	
6	53.375	0.0400	1.0	0.0	1.0	1.0	0.0	
7	53.550	0.0400	1.0	0.0	1.0	1.0	0.0	
8	53.425	0.0400	1.0	0.0	1.0	1.0	0.0	
9	53.375	0.0400	1.0	0.0	1.0	1.0	0.0	

	Pressure_switch	Oil_level	Caudal_impulses
0	1.0	1.0	1.0
1	1.0	1.0	1.0
2	1.0	1.0	1.0
3	1.0	1.0	1.0
4	1.0	1.0	1.0
5	1.0	1.0	1.0
6	1.0	1.0	1.0
7	1.0	1.0	1.0
8	1.0	1.0	1.0
9	1.0	1.0	1.0

0	1.0	1.0	1.0
1	1.0	1.0	1.0
2	1.0	1.0	1.0
3	1.0	1.0	1.0
4	1.0	1.0	1.0
5	1.0	1.0	1.0
6	1.0	1.0	1.0
7	1.0	1.0	1.0
8	1.0	1.0	1.0
9	1.0	1.0	1.0

4) Add a label column

```
#Create a new column for target variable called status, indicate the equipment has deficiencies and need to be maintained
# status = 0; system ups and running
# status = 1; system downs and needs recovering
labeled_data = data.copy()
labeled_data['status'] = 0
print(labeled_data.head(5))
```

	timestamp	TP2	TP3	H1	DV_pressure	Reservoirs	\
0	2020-02-01 00:00:00	-0.012	9.358	9.340	-0.024	9.358	
1	2020-02-01 00:00:10	-0.014	9.348	9.332	-0.022	9.348	
2	2020-02-01 00:00:19	-0.012	9.338	9.322	-0.022	9.338	
3	2020-02-01 00:00:29	-0.012	9.328	9.312	-0.022	9.328	
4	2020-02-01 00:00:39	-0.012	9.318	9.302	-0.022	9.318	

	Oil_temperature	Motor_current	COMP	DV_eletric	Towers	MPG	LPS	\
0	53.600	0.0400	1.0	0.0	1.0	1.0	0.0	
1	53.675	0.0400	1.0	0.0	1.0	1.0	0.0	
2	53.600	0.0425	1.0	0.0	1.0	1.0	0.0	
3	53.425	0.0400	1.0	0.0	1.0	1.0	0.0	
4	53.475	0.0400	1.0	0.0	1.0	1.0	0.0	

	Pressure_switch	Oil_level	Caudal_impulses	status
0	1.0	1.0	1.0	0
1	1.0	1.0	1.0	0
2	1.0	1.0	1.0	0
3	1.0	1.0	1.0	0
4	1.0	1.0	1.0	0

Next, set the status of the machine to 1 based on the time from the table below

Nr.	Start Time	End Time	Failure	Severity	Report
#1	4/18/2020 0:00	4/18/2020 23:59	Air leak	High stress	
#1	5/29/2020 23:30	5/30/2020 6:00	Air Leak	High stress	Maintenance on 30Apr at 12:00
#3	6/5/2020 10:00	6/7/2020 14:30	Air Leak	High stress	Maintenance on 8Jun at 16:00
#4	7/15/2020 14:30	7/15/2020 19:00	Air Leak	High stress	Maintenance on 16Jul at 00:00

```
def to_datetime(xs):
    result = []
    format = "%Y-%m-%d %H:%M:%S"
    for x in xs:
        result.append(pd.to_datetime(x, format = format))
    return result
```

```

failure_start_time = to_datetime(["2020-04-18 00:00:00", "2020-05-29 23:30:00", "2020-06-05 10:00:00", "2020-07-15 14:30:00"] )
failure_end_time   = to_datetime(["2020-04-18 23:59:00", "2020-05-30 06:00:00", "2020-06-07 14:30:00", "2020-07-15 19:00:00"] )

print(failure_start_time, "\n", failure_end_time[0].minute)

[Timestamp('2020-04-18 00:00:00'), Timestamp('2020-05-29 23:30:00'), Timestamp('2020-06-05 10:00:00'), Timestamp('2020-07-15 14:30:00')]
59

```

```

def in_between(x, start, end):

```

```

    start_con = x >= start
    end_con = x <= end

    inbetween_con = start_con and end_con
    if inbetween_con:
        return 1
    else:
        return 0

```

```

failure_indx = []
import numpy as np
for i, (start_time, end_time) in enumerate(zip(failure_start_time, failure_end_time)):
    mask = labeled_data['timestamp'].apply(in_between, start = start_time, end = end_time)
    indx = labeled_data.index[mask == True].tolist()
    failure_indx += indx

```

```

print(f" Found {len(failure_indx)} samples representing failure state")

```

```

    Found 29954 samples representing failure state

```

```

#Set the sample with the timestamp failed between the failure time to 1
labeled_data['status'].iloc[failure_indx] = 1

```

```

<ipython-input-16-502c3a3e2569>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
labeled_data['status'].iloc[failure_indx] = 1

```

```

#Check the number of possitive samples

```

```

print(f"We have {labeled_data['status'][labeled_data['status']==1].count()} positive samples" )

```

```

    We have 29954 positive samples

```

```

print(f"Example of Failure state \n {labeled_data[labeled_data['status']==1].head()}")

```

```

Example of Failure state
      timestamp  TP2  TP3  H1  DV_pressure  Reservoirs  \
562564 2020-04-18 00:00:01 -0.018  8.248  8.238    -0.024    8.248
562565 2020-04-18 00:00:13 -0.018  8.248  8.238    -0.024    8.248
562566 2020-04-18 00:00:24 -0.018  8.248  8.238    -0.024    8.248
562567 2020-04-18 00:00:36 -0.018  8.248  8.238    -0.024    8.248
562568 2020-04-18 00:00:49 -0.018  8.248  8.238    -0.024    8.248

      Oil_temperature  Motor_current  COMP  DV_eletric  Towers  MPG  LPS  \
562564           49.45           0.04    1.0         0.0     1.0  1.0  0.0
562565           49.45           0.04    1.0         0.0     1.0  1.0  0.0
562566           49.45           0.04    1.0         0.0     1.0  1.0  0.0
562567           49.45           0.04    0.0         0.0     0.0  0.0  0.0
562568           49.45           0.04    1.0         0.0     1.0  1.0  0.0

      Pressure_switch  Oil_level  Caudal_impulses  status
562564             1.0         1.0             1.0       1
562565             1.0         1.0             1.0       1
562566             1.0         1.0             1.0       1
562567             0.0         0.0             0.0       1
562568             1.0         1.0             1.0       1

```

5) Subsample the dataset

```
#Seperate Positive samples and Negative sample
pos_data = labeled_data[labeled_data['status'] == 1]
neg_data = labeled_data[labeled_data['status'] == 0]

#Print out the info of 2 dataset
print(f"Positive dataset\n {pos_data.info()}\n")
print(f"Negative dataset\n {neg_data.info()}\n")

<class 'pandas.core.frame.DataFrame'>
Int64Index: 29954 entries, 562564 to 1172714
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   timestamp              29954 non-null  datetime64[ns]
1   TP2                    29954 non-null  float64
2   TP3                    29954 non-null  float64
3   H1                     29954 non-null  float64
4   DV_pressure            29954 non-null  float64
5   Reservoirs             29954 non-null  float64
6   Oil_temperature        29954 non-null  float64
7   Motor_current          29954 non-null  float64
8   COMP                   29954 non-null  float64
9   DV_eletric             29954 non-null  float64
10  Towers                 29954 non-null  float64
11  MPG                    29954 non-null  float64
12  LPS                    29954 non-null  float64
13  Pressure_switch         29954 non-null  float64
14  Oil_level               29954 non-null  float64
15  Caudal_impulses        29954 non-null  float64
16  status                  29954 non-null  int64
dtypes: datetime64[ns](1), float64(15), int64(1)
memory usage: 4.1 MB
Positive dataset
None
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1486994 entries, 0 to 1516947
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   timestamp              1486994 non-null  datetime64[ns]
1   TP2                    1486994 non-null  float64
2   TP3                    1486994 non-null  float64
3   H1                     1486994 non-null  float64
4   DV_pressure            1486994 non-null  float64
5   Reservoirs             1486994 non-null  float64
6   Oil_temperature        1486994 non-null  float64
7   Motor_current          1486994 non-null  float64
8   COMP                   1486994 non-null  float64
9   DV_eletric             1486994 non-null  float64
10  Towers                 1486994 non-null  float64
11  MPG                    1486994 non-null  float64
12  LPS                    1486994 non-null  float64
13  Pressure_switch         1486994 non-null  float64
14  Oil_level               1486994 non-null  float64
15  Caudal_impulses        1486994 non-null  float64
16  status                  1486994 non-null  int64
dtypes: datetime64[ns](1), float64(15), int64(1)
memory usage: 204.2 MB
Negative dataset
None
```

As we can see, we have around 30K positive samples and 1500K negative sample. This indicates highly imbalanced dataset. Thus, we have to subsample the negative class to balance the training data. To achieve this, we will randomly sample 30K negative sample from the set of 1500K sample.

```
n_positives = int(pos_data['status'].count())
sub_neg_data = neg_data.sample(n_positives, random_state = 42)
print(f"Negative dataset after subsampling {sub_neg_data.info()}")
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29954 entries, 1306134 to 1276653
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   timestamp              29954 non-null  datetime64[ns]
1   TP2                    29954 non-null  float64
2   TP3                    29954 non-null  float64
3   H1                     29954 non-null  float64
```

```

4 DV_pressure      29954 non-null float64
5 Reservoirs      29954 non-null float64
6 Oil_temperature  29954 non-null float64
7 Motor_current    29954 non-null float64
8 COMP            29954 non-null float64
9 DV_eletric      29954 non-null float64
10 Towers         29954 non-null float64
11 MPG            29954 non-null float64
12 LPS            29954 non-null float64
13 Pressure_switch 29954 non-null float64
14 Oil_level       29954 non-null float64
15 Caudal_impulses 29954 non-null float64
16 status          29954 non-null int64
dtypes: datetime64[ns](1), float64(15), int64(1)
memory usage: 4.1 MB
Negative dataset after subsampling None

```

Now, we merge the postive set and negative set into one

```

merged_data = pd.concat([pos_data, sub_neg_data], axis = 0)
print(f"Merged dataset\n")
merged_data.info()

```

```

Merged dataset

<class 'pandas.core.frame.DataFrame'>
Int64Index: 59908 entries, 562564 to 1276653
Data columns (total 17 columns):
#   Column              Non-Null Count  Dtype
---  -
0   timestamp           59908 non-null  datetime64[ns]
1   TP2                 59908 non-null  float64
2   TP3                 59908 non-null  float64
3   H1                  59908 non-null  float64
4   DV_pressure         59908 non-null  float64
5   Reservoirs          59908 non-null  float64
6   Oil_temperature     59908 non-null  float64
7   Motor_current       59908 non-null  float64
8   COMP                59908 non-null  float64
9   DV_eletric          59908 non-null  float64
10  Towers              59908 non-null  float64
11  MPG                 59908 non-null  float64
12  LPS                 59908 non-null  float64
13  Pressure_switch     59908 non-null  float64
14  Oil_level           59908 non-null  float64
15  Caudal_impulses     59908 non-null  float64
16  status              59908 non-null  int64
dtypes: datetime64[ns](1), float64(15), int64(1)
memory usage: 8.2 MB

```

6) Find and drop outliers

```

def investigate_outliers(data, c):
    q1 = data[c].quantile(0.25)
    q3 = data[c].quantile(0.75)
    iqr = q3 - q1
    ll = q1 - 1.5*iqr
    ul = q3 + 1.5*iqr

    num_outliers = data[data[c] < ll][c].count() + data[data[c] > ul][c].count()
    if num_outliers>0:
        print(f"Found {num_outliers} outlier(s) for feature {c}")
    return {'col': c, 'n_outliers': num_outliers, 'll': ll, 'ul': ul, 'q1': q1, 'q3':q3}

print("\nDropping outliers ...\n")
clean_data = merged_data.copy()
for i in range(5):
    for c in clean_data.columns:
        if c not in ["Unnamed: 0", "timestamp"]:
            cue = investigate_outliers(clean_data, c)
            if cue["n_outliers"] > 0 and (cue["q1"]!= cue["q3"]):
                print(f"Dropping {cue['n_outliers']} from column {c}")
                clean_data = clean_data[clean_data[c]> cue["ll"]]
                clean_data = clean_data[clean_data[c]< cue["ul"]]
                print(f"{clean_data.shape[0]} samples left\n")
            elif (cue["q1"]== cue["q3"]):
                print("Skipping .. data has Q1 equals to Q3")

```

```

print(f"{clean_data.shape[0]} rows left\n")

print("\nDropping Completed ...\n")
#Recheck data
for c in clean_data.columns:
    if c not in ["Unnamed: 0", "timestamp", "COMP", 'status']:
        cue = investigate_outliers(clean_data, c)

        Found 2897 outlier(s) for feature Oil_level
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 1947 outlier(s) for feature Caudal_impulses
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 395 outlier(s) for feature LPS
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 402 outlier(s) for feature Pressure_switch
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 2897 outlier(s) for feature Oil_level
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 1947 outlier(s) for feature Caudal_impulses
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 395 outlier(s) for feature LPS
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 402 outlier(s) for feature Pressure_switch
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 2897 outlier(s) for feature Oil_level
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Found 1947 outlier(s) for feature Caudal_impulses
        Skipping .. data has Q1 equals to Q3
        59445 rows left

        Dropping Completed ...

#Investigate the columns with the binary values
binary_cols = ['LPS', 'Pressure_switch', 'Oil_level', 'Caudal_impulses']
#Ensure the the binary data is binary
clean_data[binary_cols] = clean_data[binary_cols].apply(np.round)

```

7) Summary

In summary, prior to undergoing preprocessing, the dataset:

- has unnecessary column
- has columns with wrong format

- is unlabeled
- is highly imbalanced
- has outliers

Thus, throughout the preprocessing and data cleaning phase, we performed the following tasks:

- remove unnecessary column
- format the timestamp column
- add column for target variable
- subsample to balance the dataset
- find and drop all outliers

According to its documentation, the following preprocessing steps have been conducted before publishing the data, so we do not apply it in our work:

- Data segmentation
- Normalization
- Feature Extraction

III. Exploratory Data Analysis

1) Correlation

Describing the correlation between the features, the values closer to 1 or -1 represent a stronger relation.

```
clean_data.corr().round(2)
```

```
<ipython-input-24-ef7828e8fbfe>:1: FutureWarning: The default value of numeric_only in
clean_data.corr().round(2)
```

	TP2	TP3	H1	DV_pressure	Reservoirs	Oil_temperature	Motor_curre
TP2	1.00	-0.42	-0.98	0.81	-0.42	0.74	0.87
TP3	-0.42	1.00	0.54	-0.55	1.00	-0.28	-0.17
H1	-0.98	0.54	1.00	-0.83	0.54	-0.73	-0.84
DV_pressure	0.81	-0.55	-0.83	1.00	-0.55	0.73	0.72
Reservoirs	-0.42	1.00	0.54	-0.55	1.00	-0.28	-0.17
Oil_temperature	0.74	-0.28	-0.73	0.73	-0.28	1.00	0.79
Motor_current	0.87	-0.17	-0.84	0.72	-0.17	0.79	1.00
COMP	-0.98	0.49	0.99	-0.83	0.49	-0.75	-0.84
DV_electric	0.98	-0.48	-0.99	0.83	-0.48	0.76	0.72
Towers	-0.52	0.27	0.54	-0.45	0.27	-0.41	-0.17
MPG	-0.97	0.48	0.99	-0.83	0.48	-0.75	-0.84
LPS	0.06	-0.18	-0.07	-0.05	-0.18	0.10	0.01
Pressure_switch	0.00	0.01	0.01	0.02	0.01	0.02	0.01
Oil_level	0.19	-0.14	-0.20	0.21	-0.14	0.10	0.01
Caudal_impulses	0.14	-0.13	-0.15	0.17	-0.13	0.11	0.01
status	0.83	-0.55	-0.86	0.92	-0.55	0.81	0.87


We can see that our target variable "status" has high correlation with TP2, H1, DV_pressure, Oil_temperature, Motor_current, COMP, DV_electric and MPG.

2) Visualize Correlation

Below shows a Heat map, which can be used to analyse trends, from the below heat map you can see the trends in correlation of data.


```
sns.heatmap(clean_data.corr().round(2),annot=False )
```

```
<ipython-input-25-c75868ede4c3>:1: FutureWarning: The default value of numeric_only in
sns.heatmap(clean_data.corr().round(2),annot=False )
<Axes: >
```



```
sns.pairplot(clean_data, y_vars = ['status'], plot_kws= {'alpha': 0.1})
```

```
<seaborn.axisgrid.PairGrid at 0x78d368e0d4e0>
```




3) Visualize Outliers

Drawing box plot to find outliers, I plot it on scale data so it is easier to visualize different features' range. As we can see our preprocessing function work perfectly that leaves no outliers

```
sns.set(rc={'figure.figsize':(20,8.27)})
```

```
sns.boxplot(clean_data, autorange = True)
```

<Axes: >



▼ IV. Save Data

```
clean_data.to_csv('Group_14_Clean_Data.csv')  
np.savez("Group_14_Clean_Data.npz", clean_data.to_numpy())
```